

Priyanka Gottapu

Saugus, MA | gottapupriyanka0497@gmail.com | +1 (339)-231-3729 | [Linkedin](#) | [Portfolio](#)

PROFESSIONAL SUMMARY:

- Senior Snowflake Developer / Snowflake Platform Engineer with 6+ years of experience designing, developing, modernizing, and optimizing enterprise-scale cloud-native data platforms using Snowflake, Python, SQL, AWS, Azure, Spark, and modern ELT frameworks.
- Proven expertise in Snowflake architecture, Snowpark Python, Cortex AI capabilities, Streams, Tasks, Dynamic Tables, Snowpipe, UDFs, micro-partitioning, clustering strategies, workload balancing, query optimization, RBAC, masking policies, and enterprise-grade data governance.
- Strong experience building scalable ELT/ETL pipelines using Snowflake, dbt, Fivetran, Airflow, Informatica, AWS Glue, Kafka, and Azure Data Factory supporting batch, CDC, real-time, and near real-time enterprise data processing workloads.
- Experienced in modernizing legacy enterprise systems including Workday, Salesforce, ServiceNow, Oracle, SQL Server, Teradata, and Banner into AI-ready Snowflake cloud ecosystems.
- Hands-on expertise enabling AI/ML workloads using Snowpark, Cortex, vectorized data processing, embeddings, RAG pipelines, LLM orchestration workflows, and scalable analytics architectures supporting enterprise automation and intelligent data products.
- Strong background in dimensional modeling, Data Vault 2.0, star/snowflake schema design, data lakes, data marts, and enterprise data warehousing principles.
- Experienced in cloud-native engineering across AWS and Azure environments with expertise in Terraform, CI/CD pipelines, GitHub Actions, Azure DevOps, Jenkins, infrastructure automation, API integrations, event streaming frameworks, production support, observability, disaster recovery, and enterprise security compliance.
- Proven ability to collaborate with technical and business stakeholders to translate business requirements into scalable, high-performance Snowflake data solutions.

TECHNICAL SKILLS

- ✓ **Programming Languages:** Python, SQL, JavaScript, Bash, HTML, CSS, jQuery
- ✓ **Frameworks & Libraries:** Django, Django REST Framework (DRF), FastAPI
- ✓ **Data Engineering & Snowflake:** Snowflake (Snowpipe, Streams, Tasks, Time Travel, Zero-Copy Cloning, UDFs), ETL/ELT Pipelines, Star Schema, External Functions, Snowflake Schema, Data Warehousing, Data Modeling, SnowSQL, Data Testing, Data Lineage, Metadata Management, CDC (Change Data Capture), Data Vault 2.0, Data Lakes, Medallion Architecture, Dimensional Modeling, Data Migration (Teradata, Oracle, MySQL, SQL Server → Snowflake)
- ✓ **Data Tools:** dbt, Fivetran, Snowpipe, Streams, Tasks
- ✓ **Data Integration & Transformation:** dbt Models, dbt Macros, dbt Testing, Incremental Models, API Integration, CDC Processes, Batch Processing, Real-Time Data Processing
- ✓ **Databases:** Snowflake, PostgreSQL, MySQL, SQL Server, MongoDB, Teradata

- ✓ **Cloud & AWS Services:** AWS (EC2, S3, RDS, Lambda, Glue, Redshift, Athena, IAM), Azure(ADF, Blob Storage), GCP
- ✓ **Big Data & Streaming:** Apache Kafka, Spark, Hadoop
- ✓ **SQL Expertise:** Window Functions, Query Optimization, CTEs, Semi-structured Data (ARRAY/VARIANT)
- ✓ **Data Integration & BI Tools:** Tableau, Power BI
- ✓ **Scheduling & Orchestration:** UC4 (Automatic Scheduler), Airflow
- ✓ **Web Technologies (Frontend):** HTML, CSS, JavaScript, jQuery, Vue.js
- ✓ **Tools:** Git
- ✓ **Concepts & Methodologies:** RESTful APIs, Microservices Architecture, Data Validation & Performance Tuning, RBAC (Role-Based Access Control), Agile/Scrum

PROFESSIONAL EXPERIENCE

CVS Health

Snowflake Developer

New York, USA (06/2024 to present)

- Designed and implemented enterprise-scale Snowflake cloud-native data platforms supporting healthcare analytics, operational reporting, AI-ready workloads, and governed enterprise data products across AWS environments.
- Built scalable ELT/ETL pipelines using Snowflake, Snowpark Python, PySpark, AWS Glue, Kafka, Snowpipe, Streams, Tasks, and Dynamic Tables supporting batch, CDC, real-time, and near real-time healthcare data ingestion workflows.
- Modernized legacy enterprise data structures and migrated large-scale workloads from Oracle, SQL Server, and on-premises platforms into Snowflake architectures optimized for scalability, analytics, and AI enablement.
- Developed advanced Snowflake SQL transformations, dimensional models, star schemas, Data Vault models, and optimized reporting datasets supporting enterprise business intelligence and operational analytics initiatives.
- Implemented Snowflake Cortex capabilities, vectorized processing frameworks, embeddings pipelines, and AI-ready data models supporting intelligent automation, search optimization, and RAG-based analytics use cases.
- Optimized Snowflake virtual warehouses, clustering keys, micro-partitioning strategies, caching mechanisms, and workload balancing processes significantly improving query performance and reducing cloud operational costs.
- Developed scalable data transformation frameworks using dbt, Fivetran, Python, SQL, and Airflow supporting automated data orchestration, transformation governance, and enterprise analytics delivery.

- Implemented enterprise-grade Snowflake security frameworks including RBAC, masking policies, encryption controls, secure views, audit logging, row-level security, and secure data sharing supporting HIPAA and compliance requirements.
- Integrated Snowflake with enterprise applications including Workday, Salesforce, ServiceNow, REST APIs, AWS S3, Lambda, Event Streaming frameworks, and Kafka enabling scalable cloud-native data integration ecosystems.
- Built reusable Python automation utilities and monitoring frameworks for metadata extraction, data lineage tracking, SLA monitoring, data quality validation, and pipeline observability.
- Collaborated with architects, business stakeholders, analysts, and cross-functional Agile teams to translate business requirements into scalable Snowflake data models and enterprise data engineering solutions.
- Developed CI/CD workflows using GitHub Actions, Terraform, Azure DevOps, and automated deployment pipelines supporting reliable Snowflake release management and infrastructure automation.
- Supported production environments through root cause analysis, incident resolution, query troubleshooting, warehouse tuning, disaster recovery planning, and operational support activities ensuring high platform availability.
- Conducted architecture reviews, code reviews, sprint planning sessions, technical mentoring, and engineering best-practice discussions improving development standards and enterprise delivery quality.
- Enabled enterprise AI workflows using Snowpark ML, vector search pipelines, and scalable Snowflake orchestration frameworks supporting advanced analytics and intelligent enterprise automation initiatives.

Cognizant

Snowflake Developer

Hyderabad, India (05/2020 to 06/2022)

- Led enterprise cloud migration initiatives modernizing Teradata, Oracle, SQL Server, and legacy enterprise data platforms into Snowflake cloud-native data warehouse architectures.
- Designed and developed scalable Snowflake ELT frameworks using Snowpark Python, SQL, dbt, Fivetran, AWS Glue, PySpark, Kafka, and Airflow supporting enterprise-scale batch and real-time data integration.
- Built enterprise data models including dimensional models, star schema, snowflake schema, and Data Vault architectures supporting analytics, reporting, and governed data product initiatives.
- Developed optimized Snowflake SQL queries, stored procedures, UDFs, CTEs, and transformation pipelines improving reporting performance, analytics scalability, and operational efficiency.
- Implemented Snowflake Streams, Tasks, Snowpipe, Dynamic Tables, and CDC frameworks enabling automated near real-time data ingestion and transformation workflows.
- Administered Snowflake environments including warehouse provisioning, workload balancing, clustering optimization, cost governance, RBAC implementation, masking policies, and secure access management.

- Collaborated with enterprise teams integrating Snowflake with Workday, Salesforce, ServiceNow, REST APIs, AWS services, and Kafka-based event streaming platforms.
- Developed enterprise reporting datasets, KPI dashboards, and analytical data marts supporting operational intelligence, financial reporting, customer analytics, and executive decision-making.
- Built scalable data quality validation frameworks, metadata management utilities, audit reporting pipelines, and data lineage tracking processes improving governance and compliance standards.
- Implemented CI/CD deployment pipelines using GitHub, Jenkins, Terraform, and Azure DevOps supporting automated infrastructure deployment and Snowflake release management.
- Optimized Spark and Snowflake workloads using partition strategies, clustering keys, query tuning, warehouse sizing, caching techniques, and micro-partition optimization reducing execution time and infrastructure cost.
- Worked closely with Agile/Scrum teams participating in sprint planning, backlog refinement, requirement analysis, release coordination, code reviews, and stakeholder demonstrations.
- Supported production Snowflake environments through incident management, root cause analysis, performance troubleshooting, disaster recovery support, and operational stabilization activities.
- Built Python-based automation utilities supporting ETL orchestration, API integrations, monitoring, SLA tracking, operational alerts, and enterprise pipeline observability.
- Enabled scalable cloud-native analytics ecosystems leveraging AWS, Snowflake, Spark, and enterprise integration frameworks supporting modernization and digital transformation initiatives.

Invincible Software Solutions Integrated Pvt. Ltd (ISSI)

Python Engineer

Hyderabad, India (08/2018 to 04/2020)

- Developed scalable backend systems and RESTful APIs using Django, FastAPI, and Docker, enabling seamless integration with data platforms.
- Built responsive front-end interfaces using Vue.js, HTML, CSS, and JavaScript, and designed efficient MongoDB/PostgreSQL schemas for application data.
- Developed scalable backend systems using Python and Django, following clean architecture and modular design principles.
- Designed and implemented RESTful APIs using Django REST Framework (DRF) for seamless frontend-backend communication.
- Built dynamic web applications using Django models, views, forms, and templates, ensuring high performance and maintainability.
- Developed microservices using FastAPI and integrated them with existing Django applications for improved scalability.
- Designed and optimized relational databases (PostgreSQL, MSSQL), including schema design, indexing, and query performance tuning.
- Worked with MongoDB to handle unstructured data and built efficient query patterns for high-volume applications.

- Implemented authentication, authorization, and role-based access control (RBAC) for secure application access.
- Containerized applications using Docker and deployed services on AWS (EC2, S3, RDS, Lambda) for scalable production environments.
- Built and integrated data pipelines using Python and Snowflake, including data ingestion from S3 and transformation using SQL and stored procedures.
- Supported data engineering workflows, including ETL processing, data validation, and optimization of queries in Snowflake for analytics use cases.

EDUCATION:

- Vignan's Institute of Engineering for Women | Bachelor of Science in Computer Science and Engineering (CSE) | Vishakhapatnam, India
- Rivier University | Master Of Science in Computer Information Systems (CIS) | Nashua, New Hampshire, USA

CERTIFICATIONS:

- Snowflake SnowPro Core Certification
- AWS Certified Cloud Practitioner CLF-C02
- Data Analytics by Cisco Networking